

ATHARVA JIBHAKATE

atharvajibhakate@gmail.com | +91-7020413721 | [linkedin.com/in/atharva-jibhakate](https://www.linkedin.com/in/atharva-jibhakate) | [GitHub](https://github.com)

SUMMARY

Senior Data Engineer with 4+ years of experience architecting enterprise data platforms on Databricks and AWS, processing 500GB–10TB+ daily across adtech and healthcare domains. Expertise in lakehouse design (Delta Lake, Apache Iceberg), medallion architecture, CDC pipelines, and Unity Catalog data governance. Pioneered LLM-powered automation via Databricks Model Serving, eliminating manual metadata overhead across 300+ Delta tables. Delivered measurable impact — 58% ETL runtime reduction, 99%+ SLA adherence, and 60% operational effort savings — across production-grade, high-throughput ELT pipelines.

EXPERIENCE

Data Engineer | Publicis Groupe | Pune, India

Feb 2025 – Present

Product: GrowthOS - Insights and Analyze

- Architected Databricks-based lakehouse pipelines (PySpark, Delta Lake) ingesting data from 65+ ad platforms (Google, YouTube, TikTok, Snapchat) within Bronze–Silver–Gold architecture, processing 500GB+ daily, and improving data freshness from T+1 to near real-time ingestion.
- Modernized data platform using Delta Live Tables (DLT) with declarative data quality expectations, reducing pipeline failures by 30% and enabling lineage tracking across ingestion layers.
- Automated outbound data delivery to client Snowflake via Delta Sharing and Redis workspaces via API-driven workflows, replacing manual exports and reducing operational effort by 60% and delivery latency by 40%.
- Deployed an LLM-powered metadata pipeline using Databricks Serving to auto-generate table and column metadata for 300+ Delta tables across 65+ data sources, surfacing annotations in Unity Catalog Explorer, and enabling AI-based governance tagging consumed by downstream pipelines and access-control workflows.
- Designed a Digital Performance Metrics (DPM) framework standardizing 100+ cross-platform KPIs into a unified reporting layer, improving metric consistency by 35% and eliminating cross-dashboard discrepancies for product stakeholders.
- Orchestrated end-to-end multi-tenant data workflows on Amazon MWAA in collaboration with the DevOps team, implementing modular Airflow DAGs coordinating layered PySpark jobs across medallion transformations and achieving 99%+ SLA adherence for client-facing campaign delivery.
- Engineered scalable Databricks job frameworks integrating PaaS APIs for inferred schema mapping and synchronization, reducing new data-source onboarding time by 40%.

Data Engineer | UST Global | Pune, India

May 2022 – Feb 2025

Client: Blue Cross and Blue Shield Association

- Led end-to-end architecture and delivery of scalable AWS-based ETL batch processing pipelines using Glue, Spark, Lambda, and EC2, integrating healthcare insurance data from S3, PostgreSQL, and Elasticsearch into analytics-ready layers.
- Reduced ETL runtime from 60 to 25 minutes (58% improvement) by eliminating data skew, implementing broadcast joins, repartitioning strategies, and advanced Spark performance tuning techniques—cutting compute costs by 40%.
- Engineered a high-performance CDC framework using Apache Iceberg and Spark SQL, improving incremental processing efficiency by 30% and enabling reliable schema evolution across the data warehousing layers.
- Processed and indexed 10TB+ of daily data into Elasticsearch using Glue jobs, enabling near-real-time analytical access for downstream reporting systems.
- Spearheaded system-wide performance optimization initiatives using CloudWatch and X-Ray, improving pipeline execution speed by 50% and enhancing production stability.
- Analysed and optimized Redshift workloads using Spectrum and concurrency scaling, improving query performance by 20% while supporting 10,000+ concurrent queries during peak loads.

TECHNICAL SKILLS

Design & Architecture: Lakehouse Architecture, Medallion (Bronze–Silver–Gold), Data Warehousing, Data Modeling (Star/Snowflake Schema), Batch Processing, CDC, Distributed Systems, Schema Evolution, Data Governance, Data Quality Frameworks

Big Data & Processing: ETL/ELT, Apache Spark, PySpark, Databricks, Delta Lake, Delta Live Tables (DLT), Apache Iceberg, Pandas

Cloud & Data Platforms: AWS (Glue, Lambda, EMR, S3, EC2, Redshift, Athena, SNS), Snowflake, Elasticsearch, Redis

Orchestration & DevOps: Apache Airflow (MWAA), AWS Step Functions, Git, CI/CD Pipelines, Docker, IaC (Terraform)

Languages: Python, SQL, Java, Shell Scripting

EDUCATION

MIT World Peace University | Pune, Maharashtra, India | Jul 2018 – Jun 2022

Bachelor of Technology in Computer Science Engineering | 9.1/10 CGPA

CERTIFICATIONS

Databricks Certified Associate Data Engineer | Databricks | Jul 2025 – Jul 2027

AWS Certified Associate Data Engineer | Amazon Web Services | Aug 2023 – Aug 2026